

Machine Learning Internship Report

Submitted by: Yatharth Kumar Saxena

Objective:

To build a classification model using Random Forest to predict the Target variable from customer data. The goal is to compare performance between raw and preprocessed datasets using two different train-test splits: **70:30** and **80:20**.

Step-by-Step Explanation

1. Importing Required Libraries

The project started by importing essential Python libraries:

- pandas – for data handling
- numpy – for numerical computations
- scikit-learn – for model building, data preprocessing, and evaluation

These libraries enabled efficient data loading, transformation, model training, and performance assessment.

2. Loading and Preparing the Dataset

The dataset was loaded from a CSV file named:

Yatharth Kumar Saxena - ml_preprocessing_dataset_1000.csv

A .copy() of the dataset was stored in two versions:

- df_raw → to retain the raw dataset structure (no imputations)
- df_clean → to apply full preprocessing and transformations

The columns were stripped of extra whitespace to ensure consistency.

3. Initial Transformation – Target Encoding

The Target column was found to be of type object. To allow model training:

- It was label-encoded in both df_raw and df_clean.

4. Raw vs Cleaned Dataset Processing

Raw Dataset (df_raw)

- No missing value handling
- All categorical (string) features except Legacy_Customer_ID and Target were **label encoded directly**

Cleaned Dataset (df_clean)

- Categorical missing values were filled with **mode**
- Numerical missing values were filled with **mean**

- After imputation, all object-type columns were label-encoded

5. Feature-Target Separation

The dataset was split into:

- $x \rightarrow$ All columns except Target and Legacy_Customer_ID
- $y \rightarrow$ Only the Target column

This separation was done uniformly for both raw and cleaned datasets.

6. Model Evaluation Pipeline

A reusable function `evaluate_model()` was created to:

- Split data into **training**, **validation**, and **testing** sets
- Apply `StandardScaler` if `preprocess=True`
- Train a **Random Forest Classifier**
- Print:
 - Accuracy
 - Classification Report (Precision, Recall, F1)
 - Confusion Matrix

Both 70:30 and 80:20 train-test splits were tested with and without preprocessing.

Results

Split: 70:30 | Preprocessed: (Raw Data)

- **Accuracy:** 59.13%
- **F1-Score (Class 0 / Class 1):** 0.72 / 0.22
- **Macro Avg F1:** 0.47
- **Confusion Matrix:**

```
[[161 27]
 [ 96 17]]
```

Split: 70:30 | Preprocessed: (Cleaned Data)

- **Accuracy:** 57.80%
- **F1-Score (Class 0 / Class 1):** 0.72 / 0.17
- **Macro Avg F1:** ~0.45
- **Confusion Matrix:**

```
[[110 18]
 [ 63  9]]
```

Split: 80:20 | Preprocessed: (Raw Data)

- Accuracy: 61.5%
- F1-Score (Class 0 / Class 1): 0.75 / 0.21
- Macro Avg F1: 0.48
- Confusion Matrix: $\begin{bmatrix} 113 & 15 \\ 62 & 10 \end{bmatrix}$

Split: 80:20 | Preprocessed: (Cleaned Data)

- Accuracy: 61.5%
- F1-Score (Class 0 / Class 1): 0.75 / 0.19
- Macro Avg F1: 0.47
- Confusion Matrix: $\begin{bmatrix} 114 & 14 \\ 63 & 9 \end{bmatrix}$

Observations

- The model consistently performed **better on Class 0**, indicating a possible **class imbalance**.
- Preprocessing improved consistency and structure, but did **not significantly improve accuracy or recall** for Class 1.
- Across both splits, **Class 1 recall** remained low (~14% or less), limiting overall F1 score and macro average.
- Standard scaling and missing value treatment did not change the outcome drastically due to Random Forest's tree-based nature.

Conclusion

This internship project provided hands-on experience in building an end-to-end machine learning pipeline:

- From **data ingestion** and **cleaning**
- To **model training, evaluation, and comparison** across datasets

Although the Random Forest Classifier achieved only ~61% accuracy, the project underlines the importance of:

- **Data quality**
- **Class balancing**
- **Model choice**
- **Hyperparameter tuning**

Future enhancements may include:

- SMOTE / undersampling for **balancing**
- Trying **XGBoost, SVM, or Logistic Regression**
- Applying **GridSearchCV** for hyperparameter optimization